

Label Words are Anchors: An Information Flow Perspective for Understanding In-Context Learning

@article{wen2023hard,
title={Label Words are Anchors:
An Information Flow Perspective for Understanding In-Context Learning},
author={Wang, Lean and Li, Lei and Dai, Damai and Chen, Deli and Zhou, Hao and Meng,
Fandong and Zhou, Jie and Sun, Xu},
journal={2305.14160},
year={2023}
Conference={ENNLP}
Cites={2}
Best Paper Award

Background and Motivation

LLMs如何通过ICL来提升它们理解和生成文本的能力。LLMs可以在给定适当上下文的情况下生成有意义和相关的文本。这些模型通常需要大量的数据来训练，而且不总是能够充分理解上下文中提供的信息是如何影响输出的。

1. **资源效率**：希望找到更有效率的方法来利用上下文信息，减少资源消耗。
2. **上下文利用**：当前LLMs对上下文信息的利用还不够充分。
3. **学习过程的优化**：模型的学习过程有改进的空间，特别是在如何从示例中学习和提取有用信息的方面。

Information Flow with Labels as Anchors

\mathcal{H}_1 : In shallow layers, label words gather the information of demonstrations to form semantic representations for deeper layers.

\mathcal{H}_2 : In deep layers, the model extracts the information from label words to form the final prediction.




Figure 1: Visualization of the information flow in a GPT model performing ICL. The line depth reflects the significance of the information flow from the right word to

Method

作者利用显著性分数来揭示GPT模型中标记之间注意力交互的内在模式

$$I_l = \left| \sum_h A_{h,l}^\top \frac{\partial \mathcal{L}(x)}{\partial A_{h,l}} \right|.$$

S_{wp} , the mean significance of information flow from the text part to label words:

$$S_{wp} = \frac{\sum_{(i,j) \in C_{wp}} I_l(i,j)}{|C_{wp}|}, \quad (2)$$

$$C_{wp} = \{(p_k, j) : k \in [1, C], j < p_k\}.$$

S_{pq} , the mean significance of information flow from label words to the target position:


$$S_{pq} = \frac{\sum_{(i,j) \in C_{pq}} I_l(i,j)}{|C_{pq}|}, \quad (3)$$

$$C_{pq} = \{(q, p_k) : k \in [1, C]\}.$$


S_{ww} , the mean significance of the information flow amongst all words, excluding influences represented by S_{wp} and S_{pq} :

$$S_{ww} = \frac{\sum_{(i,j) \in C_{ww}} I_l(i,j)}{|C_{ww}|}, \quad (4)$$

$$C_{ww} = \{(i,j) : j < i\} - C_{wp} - C_{pq}.$$



(a) Results on the SST-2 dataset



(b) Results on the AGNews dataset

这三个指标用来评估模型中不同类型信息流动的强度和重要性。 S_{wp} 评估了输入文本到标签词的信息聚合程度, S_{pq} 评估了标签词对最终决策的影响程度, 而 S_{ww} 提供了一个除去上述两种特定信息流动的全局信息流动的基准

Proposed Hypothesis Based on this, we propose the hypothesis that **label words function as anchors** in the ICL information flow. In shallow layers, **label words gather information** from demonstration examples to form semantic representations for deeper layers, while in deep layers, the model **extracts the information from label words** to form the final prediction. Figure 2 gives an illustration for our hypothesis.

Method

Shallow Layers: Information Aggregation

验证信息聚合在ICL中的角色，特别是在模型浅层中标签词的作用。目的是展示当阻止了标签词从文本中获取信息时，模型行为的变化。

阻断信息流:通过操纵注意力矩阵 A , 阻断到标签词的信息流。将注意力矩阵 A 的第 l 层中, 标签词 p 位置之前的所有位置 i 的注意力值设置为0, $A(p,i)$ 对于所有 $i < p$ 。在第 l 层中, 标签词无法从先前的文本中聚合信息

当在模型的前5层隔离标签词时, 模型输出的一致性显著下降, 这表明**浅层的信息聚合**对模型预测非常重要。相比之下, 隔离后5层的标签词或随机隔离非标签词对模型预测的影响较小。这进一步证实了标签词在浅层中作为信息聚合锚点的作用。

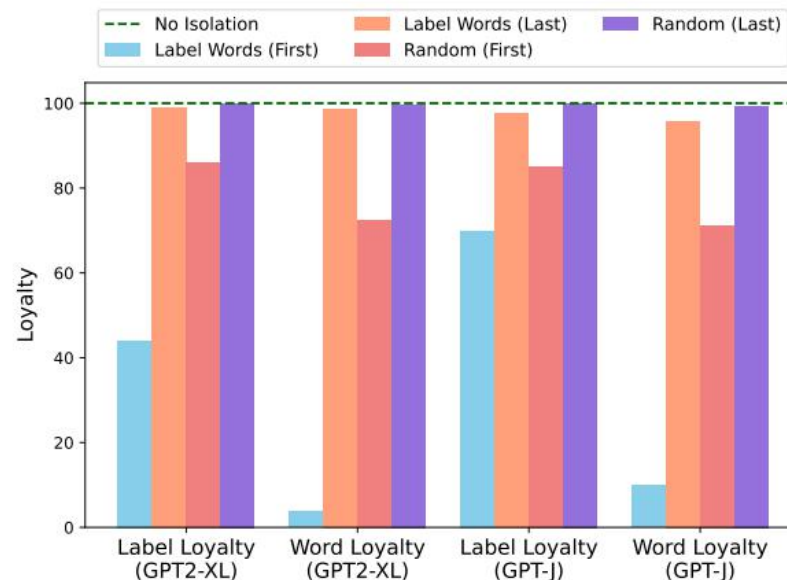


Figure 4: The impact of isolating label words versus randomly isolating non-label words within the first or last 5 layers. Isolating label words within the first 5 layers exerts the most substantial impact, highlighting the importance of shallow-layer information aggregation via label words.

Method

Deep Layers: Information Extraction


验证信息聚合在ICL中的角色，特别是在模型浅层中标签词的作用。目的是展示当阻止了标签词从文本中获取信息时，模型行为的变化。

$(A_l(q, p_1), \dots, A_l(q, p_C))$ 对最后结果的影响


We utilize the AUC-ROC score to quantify the correlation between $A_l(q, p_i)$ and model prediction, which we denote as AUCROC_l for the l -th layer.

$$R_l = \frac{\sum_{i=1}^l (\text{AUCROC}_i - 0.5)}{\sum_{i=1}^N (\text{AUCROC}_i - 0.5)}.$$

量化直到第 l 层为止所有层对模型预测的累积贡献



(a) GPT2-XL (total 48 layers).



(b) GPT-J (total 28 layers).

Anchor Re-weighting:提高ICL的准确性

$$\begin{aligned} \Pr_f(Y = i|X = x) &\approx A(q, p_i) \\ &= \frac{\exp(\mathbf{q}_q \mathbf{k}_{p_i}^T / \sqrt{d})}{\sum_{j=1}^N \exp(\mathbf{q}_q \mathbf{k}_j^T / \sqrt{d})}. \end{aligned} \quad (6)$$

By setting $\mathbf{q}_q / \sqrt{d} = \hat{\mathbf{x}}$ and $\mathbf{k}_{p_i} - \mathbf{k}_{p_C} = \beta_i$, we deduce:

$$\log \frac{\Pr_f(Y = i|X = x)}{\Pr_f(Y = C|X = x)} = \beta_i^T \hat{\mathbf{x}}. \quad (7)$$

This approximates a logistic regression model where:

$$\log \frac{\Pr_f(Y = i|X = x)}{\Pr_f(Y = C|X = x)} = \beta_0^i + \beta_i^T \mathbf{x}. \quad (8)$$

$$\hat{A}(q, p_i) = \exp(\beta_0^i) A(q, p_i) \quad (9)$$

To train the re-weighting vector $\beta = \{\beta_0^i\}$, we utilize an auxiliary training set $(\mathbf{X}_{train}, \mathbf{Y}_{train})$. Here, we perform ICL with normal demonstrations and optimize β with respect to the classification loss \mathcal{L} on $(\mathbf{X}_{train}, \mathbf{Y}_{train})$:

$$\beta^* = \arg \min_{\beta} \mathcal{L}(\mathbf{X}_{train}, \mathbf{Y}_{train}). \quad (10)$$

Method	SST-2	TREC	AGNews	EmoC	Average
Vanilla In-Context Learning (1-shot per class)	61.28	57.56	73.32	15.44	51.90
Vanilla In-Context Learning (5-shot per class)	64.75	60.40	52.52	9.80	46.87
Anchor Re-weighting (1-shot per class)	90.07	60.92	81.94	41.64	68.64

Table 1: The effect after adding parameter β_0^i . For AGNews, due to the length limit, we only use three demonstrations per class. Our Anchor Re-weighting method achieves the best performance overall tasks.

Method

Applications of Our Anchor-Based Understanding

Anchor-Only Context Compression: , 模型推理时会生成一系列隐藏状态, 每个状态对应于输入文本中的一个词。通常情况下, 模型在生成文本或推理时, 会考虑所有隐藏状态。

只考虑与“锚点”相关的隐藏状态, 而不是全部状态。在进行推理时只需要处理这些锚点对应的状态, 加快推理速度。

pendent of subsequent words. This allows for the calculation and caching of the label word hidden states $\mathbf{H} = \{\{\mathbf{h}_l^i\}_{i=1}^C\}_{l=1}^N$ (\mathbf{h}_l^i is the l -th layer's hidden state of the i -th label word in the demonstration). By concatenating $\mathbf{h}_l^1, \dots, \mathbf{h}_l^C$ at the front in each layer during inference, instead of using the full demonstration, we can speed up inference.

Text_{anchor}: This method concatenates the formatting and label text with the input, as opposed to concatenating the hidden states at each layer.

Hidden_{random}: This approach concatenates the hidden states of formatting and randomly selected non-label words (equal in number to Hidden_{anchor}).

Hidden_{random-top}: To establish a stronger baseline, we randomly select 20 sets of non-label words in Hidden_{random} and report the one with the highest label loyalty.

Method	Label Loyalty	Word Loyalty	Acc.
ICL (GPT2-XL)	100.00	100.00	51.90
Text _{anchor}	51.05	36.65	38.77
Hidden _{random}	48.96	5.59	39.96
Hidden _{random-top}	57.52	4.49	41.72
Hidden _{anchor}	79.47	62.17	45.04
ICL (GPT-J)	100.00	100.00	56.82
Text _{anchor}	53.45	43.85	40.83
Hidden _{random}	49.03	2.16	31.51
Hidden _{random-top}	71.10	11.36	52.34
Hidden _{anchor}	89.06	75.04	55.59


Method

Anchor Distances for Error Diagnosis


用于诊断ICL错误的方法，是通过检查注意力模块中对应于标签词的关键向量之间的距离

$$\text{Confusion}_{ij}^{\text{pred}} = \frac{\|\hat{\mathbf{k}}_{\mathbf{p}_i} - \hat{\mathbf{k}}_{\mathbf{p}_j}\|}{\max_{s \neq t} \|\hat{\mathbf{k}}_{\mathbf{p}_s} - \hat{\mathbf{k}}_{\mathbf{p}_t}\|}, \quad (11)$$

用了类似主成分分析（PCA）的方法从键向量中提取与最大变异量对应的成分



(a) Confusion matrix of $\text{Confusion}_{ij}^{\text{pred}}$.



(b) Confusion matrix of Confusion_{ij} .



When Do Prompting and Prefix-Tuning Work? A Theory of Capabilities and Limitations

@article{
title={When Do Prompting and Prefix-Tuning Work?
A Theory of Capabilities and Limitations},
Conference={ICLR 2024}
Cites={0}
6 6 6 8

Background and Motivation

fine-tuning, soft prompt and prefix-tuning
motivations, abilities and limitations

1. **Soft prompting and prefix-tuning** are motivated by the **embedding space** being **larger** than the **token space**. However, can a transformer utilize the additional capacity? We show that with a careful choice of transformer weights, controlling a single embedding can generate any
 2. **Since prefix-tuning is more expressive than prompting, is it as expressive as full fine-tuning?** Despite the expressiveness of continuous space, prefix-tuning has structural limitations. A prefix cannot change the relative attention over the content tokens and can only bias the output of the attention block in a constant direction. In contrast, full fine-tuning can learn new attention patterns and arbitrarily modify attention block outputs, making it strictly more powerful.
 3. **If context-based fine-tuning methods suffer from such structural limitations, how come they have high empirical performance?** We show that the prefix-induced bias can steer the model towards a pretraining task. Prefix-tuning can also combine skills picked up during pretraining to solve some new tasks similar to pretraining tasks. However, it cannot learn a completely new task. This is not simply because of the small number of learnable parameters: fine-tuning the same number of parameters can be sufficient to learn the novel task. Hence, context-based fine-tuning can elicit or combine pretrained model skills but cannot learn completely new behaviors.
-

3 SOFT PROMPTING HAS MORE CAPACITY THAN PROMPTING

Soft prompt和prefix turning是由embedding space的广度激发的，这个空间比单个token的可能完成项的空间要大。其成功的原因，主要归因于embedding space的容量大于固定token空间的容量。嵌入空间是不可数无限的，相比之下，token空间是有限的。

- 无条件生成 (unconditional generation) 指的是没有输入的情况下，系统token生成序列的能力。例如，使用一个单一的系统token S_1 生成序列 $(Y_1, \dots, Y_N) = f_{S_1}$ 。
- 如果使用确定性自回归函数，一个系统token S_1 只能生成 V 种不同的序列，因为第一个token决定了序列的其余部分。
- 但是，如果我们使用一个由实数向量索引的虚拟单一系统token s_1 ，理论上可以生成所有 V^N 可能的输出序列，因为实数向量是无限的。

V 是词汇表的大小
 N 代表生成序列的长度

有一个简单的自回归模型，它的词汇表大小是 V ，其中包含了如“apple”，“banana”，“cherry”等词汇。如果我们只改变序列的第一个词（第一个系统token S_1 ），并且这个词从词汇表中选取，那么整个序列的变化只能基于这个第一个词的选择。如果模型决定在“apple”之后总是跟着“is red”，那么每次选择“apple”作为第一个词时，生成的序列就总会是“apple is red”。既然我们只能从 V 个词汇中选择第一个词，模型最多只能生成 V 种不同的序列。

Method

这部分的结论是soft prompt和prefix turning具有比传统prompt更大的表现力。

尽管可以用virtual token完全决定从用户输入到模型响应的映射，这可能给人一种错觉，认为soft prompt的能力与full fine-tuning一样强大。

但是。soft prompt和full fine-tuning存在结构上的限制，它们无法促进学习一个全新任务的能力。接下来的部分将阐明差异。

Theorem 1 (Exponential unconditional generation capacity of a single virtual token). *For any $V, N > 0$, there exists a transformer with vocabulary size V , context size N , embedding size $d_e = N$, one attention layer with two heads and a three-layer MLP such that it generates any token sequence $(Y_1, \dots, Y_N) \in \{1, \dots, V\}^N$ when conditioned on the single virtual token $s_1 = ((Y_1 - 1)/V, \dots, (Y_N - 1)/V)$.*

from X_1 to Y_1 , but S_1 can take on only V values: $|\{f_{s_1} : S_1 \in 1, \dots, V\}| = V < V^V$. Hence, tokens cannot be used to specify an arbitrary map from user input to model output. However, a single virtual token can specify any of the V^V maps, i.e., there exists a transformer $f_{s_1}(X_2)$ for which there is a surjective map from $\{f_{s_1} : s_1 \in \mathbb{R}^{d_e}\}$ to $\{1, \dots, V\}^{\{1, \dots, V\}}$.

定理1说明了使用单一virtual token的transformer模型在理论上具有强大的生成能力，即使在条件生成的场景下，这个能力仍然非常强大，能够映射出复杂的输入到输出序列的关系。

Theorem 2 (Conditional generation capacity for a single virtual token ($n_X = n_Y = 1$)). *For any $V > 0$, there exists a transformer with vocabulary size V , context size $N = 2$, embedding size $d_e = V$, one attention layer with two heads and a three-layer MLP that reproduces any map $m: [1, \dots, V] \rightarrow [1, \dots, V]$ from a user input token to a model response token when conditioned on a single virtual token $s_1 = (m(1)/V, \dots, m(V)/V)$. That is, by selecting s_1 we control the model response to any user input.*

定理2扩展了定理1的概念。不仅是对于单个响应token，而且对于更长的响应序列，通过增加soft prompt的长度来增加用户输入的长度，soft prompt显示出更大的表现力。

While full fine-tuning can alter the attention pattern of an attention head, prefix-tuning cannot.

Recall the attention A_{ij} position i gives to position j for a trained transformer (Equation (1)):

$$A_{ij} = \frac{\exp\left(T/\sqrt{k} \mathbf{x}_i^\top \mathbf{W}_Q^\top \mathbf{W}_K \mathbf{x}_j\right)}{\sum_{r=1}^p \exp\left(T/\sqrt{k} \mathbf{x}_i^\top \mathbf{W}_Q^\top \mathbf{W}_K \mathbf{x}_r\right)} = \frac{\exp\left(T/\sqrt{k} \mathbf{x}_i^\top \mathbf{H} \mathbf{x}_j\right)}{\sum_{r=1}^p \exp\left(T/\sqrt{k} \mathbf{x}_i^\top \mathbf{H} \mathbf{x}_r\right)}, \quad (5)$$

where $\mathbf{W}_Q^\top \mathbf{W}_K = \mathbf{H}$. Full fine-tuning can enact arbitrary changes to \mathbf{W}_Q and \mathbf{W}_K and hence, assuming the input does not change (e.g., at the first attention layer), we get the following attention:

$$A_{ij}^{\text{ft}} = \frac{\exp\left(T/\sqrt{k} \mathbf{x}_i^\top \mathbf{H} \mathbf{x}_j + T/\sqrt{k} \mathbf{x}_i^\top \Delta \mathbf{H} \mathbf{x}_j\right)}{\sum_{r=1}^p \exp\left(T/\sqrt{k} \mathbf{x}_i^\top \mathbf{H} \mathbf{x}_r + T/\sqrt{k} \mathbf{x}_i^\top \Delta \mathbf{H} \mathbf{x}_r\right)},$$

where the changes to \mathbf{W}_Q and \mathbf{W}_K are folded into $\Delta \mathbf{H}$. It is clear that by varying $\Delta \mathbf{H}$ full fine-tuning can change the attention patterns arbitrarily. However, let us see how is attention affected by


$$A_{i0}^{\text{pt}} = \frac{\exp\left(T/\sqrt{k} \mathbf{x}_i^\top \mathbf{H} \mathbf{s}_1\right)}{\exp\left(\frac{T}{\sqrt{k}} \mathbf{x}_i^\top \mathbf{H} \mathbf{s}_1\right) + \sum_{r=1}^p \exp\left(\frac{T}{\sqrt{k}} \mathbf{x}_i^\top \mathbf{H} \mathbf{x}_r\right)}, \quad A_{ij}^{\text{pt}} = \frac{\exp\left(T/\sqrt{k} \mathbf{x}_i^\top \mathbf{H} \mathbf{x}_j\right)}{\exp\left(\frac{T}{\sqrt{k}} \mathbf{x}_i^\top \mathbf{H} \mathbf{s}_1\right) + \sum_{r=1}^p \exp\left(\frac{T}{\sqrt{k}} \mathbf{x}_i^\top \mathbf{H} \mathbf{x}_r\right)} \quad \text{for } j \geq 1.$$

The numerator of A_{ij}^{pt} is the same as in Equation (5), i.e., the prefix does not affect it. It only adds the term $\exp(T/\sqrt{k} \mathbf{x}_i^\top \mathbf{H} \mathbf{s}_1)$ to the denominator. Therefore, the attention position i gives to the content positions $j \geq 1$ is simply scaled down by the attention it now gives to the prefix. If *tomato* attends the most to *salad* in a particular context, no prefix can change that. This becomes evident by rewriting A_{ij}^{pt} as the attention of the pretrained model scaled by the attention “stolen” by the prefix:

前缀调整相对于完整微调，它只能在模型现有的注意力框架内引入偏差，而不能创建新的注意力模式，只能在原有的注意力模式上加权或减权

Method

4 PREFIX-TUNING CAN ONLY BIAS THE OUTPUT OF AN ATTENTION HEAD



Longer prefixes **define larger subspaces** for the bias but are **not fully utilized in practice**

有更长的前缀和偏置空间：更长的前缀会定义更大的偏置空间，但实际上并未完全利用。偏置空间并不是被前缀完全占据的。尽管理论上前缀可以定义一个很大的空间，实际中前缀的注意力分布并不依赖于输入内容，它们主要作为偏置存在

Prefix-tuning only adds a bias to the attention block output. Let us see how this attention scaling down affects the **output of the attention block**. Following Equation (2), the output at position i for the **pretrained (t_i)**, the **fully fine-tuned (t_i^{ft})** and the **prefix-tuned (t_i^{pt})** models are as follows:⁵

$$\begin{aligned}
 t_i &= \sum_{j=1}^p A_{ij} W_V x_j, & t_i^{\text{ft}} &= \sum_{j=1}^p A_{ij}^{\text{ft}} (W_V + \Delta W_V) x_j, \\
 t_i^{\text{pt}} &= A_{i0}^{\text{pt}} W_V s_1 + \sum_{j=1}^p A_{ij}^{\text{pt}} W_V x_j \stackrel{(6)}{=} A_{i0}^{\text{pt}} W_V s_1 + \sum_{j=1}^p A_{ij} (1 - A_{i0}^{\text{pt}}) W_V x_j = \underbrace{A_{i0}^{\text{pt}} W_V s_1}_{\text{bias}} + (1 - A_{i0}^{\text{pt}}) t_i.
 \end{aligned}
 \tag{7}$$

Prefix-tuning **cannot** learn a new task requiring a different attention pattern, prefix-tuning indeed cannot learn a new task if it requires **new attention patterns**.




Table 1: A transformer pretrained on sorting in ascending order cannot be prefix-tuned to sort in descending order. 10 random seeds.

	Ascending	Descending
Pretrain on asc.	$91 \pm 5\%$	$0 \pm 0\%$
Full fine-tune on desc.	$0 \pm 0\%$	$85 \pm 5\%$
Prefix-tune on desc.	$0 \pm 0\%$	$0 \pm 0\%$

Figure 1: Attention patterns of a small transformer pretrained on sorting in ascending order. The model is given the prefix S and user input X and generates Y autoregressively. We have highlighted the attention when the first response Y_1 is being generated. Full fine-tuning sorts in descending order but prefix-tuning cannot as it cannot update the learned attention. Note how the relative attention of X to X in the left and right plots is exactly the same: the prefix cannot change the attention pattern for the same inputs. The relative attention of X to X in the center plot is very different because full fine-tuning can arbitrarily change W_Q and W_K .

Prefix-tuning can **elicit a skill** from the pretrained model.

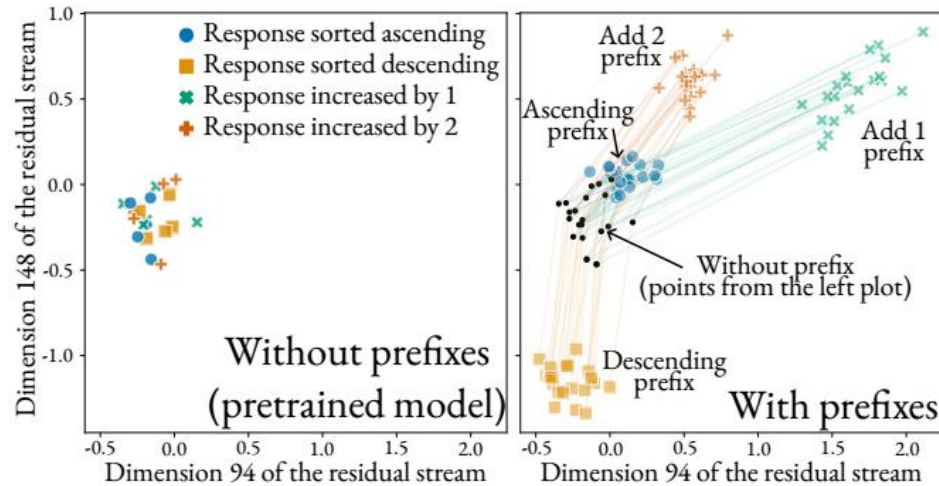


Table 2: A transformer pretrained on several tasks can be prefix-tuned for one of them. 10 random seeds.

Accuracy on:	\nearrow	\searrow	+1	+2
Pretrained	25 \pm 13%	25 \pm 12%	24 \pm 11%	22 \pm 7%
Prefix-tune on \nearrow	95 \pm 2%	0 \pm 0%	0 \pm 0%	0 \pm 0%
Prefix-tune on \searrow	0 \pm 0%	90 \pm 3%	1 \pm 1%	1 \pm 1%
Prefix-tune on +1	0 \pm 0%	1 \pm 3%	95 \pm 6%	0 \pm 1%
Prefix-tune on +2	0 \pm 0%	0 \pm 0%	1 \pm 2%	98 \pm 5%

Table 3: Prefix tuning can learn a new task requiring only pretraining skills (\nearrow +1) but cannot learn a completely new task (\mathbb{H}). Average accuracy over 3 seeds.

Accuracy on:	\nearrow	\searrow	+1	+2	\nearrow +1	\mathbb{H}
Pretrained	17%	23%	34%	25%	0%	0%
Prefix-tune on \nearrow	100%	0%	0%	0%	0%	0%
Prefix-tune on \searrow	0%	100%	0%	0%	0%	0%
Prefix-tune on +1	0%	0%	100%	0%	0%	0%
Prefix-tune on +2	0%	0%	0%	100%	0%	0%
Prefix-tune on \nearrow +1	0%	0%	0%	0%	93%	0%
Prefix-tune on \mathbb{H}	0%	0%	0%	0%	0%	1%

Hard prompts made easy: Gradient-based discrete optimization for prompt tuning and discovery

@article{wen2023hard,
title={Hard prompts made easy: Gradient-based discrete optimization for prompt tuning and discovery},
author={Wen, Yuxin and Jain, Neel and Kirchenbauer, John and Goldblum},
journal={arXiv:2302.03668},
year={2023}
Conference={Nips}
Cites={47}}

Background and Motivation

黑盒prompt 优化

Hard prompts can be automatically discovered that are effective in tuning LMs for **classification**.

Algorithm 1 Hard Prompts made EaZy: PEZ Algorithm

Input: Model θ , vocabulary embedding $\mathbf{E}^{|V|}$, projection function Proj , broadcast function \mathcal{B} , optimization steps T , learning rate γ , Dataset D

Sampled from real embeddings:

$\mathbf{P} = [\mathbf{e}_i, \dots, \mathbf{e}_M] \sim \mathbf{E}^{|V|}$

for $1, \dots, T$ **do**

 Retrieve current mini-batch $(X, Y) \subseteq D$.

 Forward Projection:

$\mathbf{P}' = \text{Proj}_{\mathbf{E}}(\mathbf{P})$

 Calculate the gradient w.r.t. the *projected* embedding:

$g = \nabla_{\mathbf{P}'} \mathcal{L}_{\text{task}}(\mathcal{B}(\mathbf{P}', X_i), Y_i, \theta)$

 Apply the gradient on the *continuous* embedding:

$\mathbf{P} = \mathbf{P} - \gamma g$

end for

Final Projection:

$\mathbf{P} = \text{Proj}_{\mathbf{E}}[\mathbf{P}]$

return \mathbf{P}

$$\mathcal{L} = (1 - \lambda_{\text{fluency}}) \mathcal{L}_{\text{task}} + \lambda_{\text{fluency}} \mathcal{L}_{\text{fluency}}.$$

Method	GPT-2 Large (755M, Source)	GPT-2 XL (1.3B)	T5-LM-XL (3B)	OPT (2.7B)	OPT (6.7B)
Empty _{Template}	80.84	73.85	52.75	72.48	58.72
AutoPrompts _{SGD}	87.56 \pm 0.35	78.19 \pm 2.68	56.01 \pm 1.67	73.69 \pm 1.63	65.28 \pm 1.75
FluentPrompt	88.33 \pm 0.35	78.53 \pm 2.82	55.64 \pm 0.59	70.39 \pm 2.08	61.74 \pm 1.25
PEZ _{No Fluency} (Ours)	88.12 \pm 0.15	77.8 \pm 3.45	61.12 \pm 2.94	76.93 \pm 1.29	71.72 \pm 3.16
PEZ _{Fluency} (Ours)	88.05 \pm 0.55	79.72 \pm 3.26	63.30 \pm 2.30	77.18 \pm 3.82	72.39 \pm 1.82